

Mathematical Physics

taught by Michael Aizenman

February 4, 2016

1	Introduction to statistical mechanics	5
1	The equivalence principle	5
1.1	Configurations and ensembles	5
1.2	Equivalence principle: first proof	6
1.3	Second proof	8
1.4	Legendre transform	9

Introduction

Notes from Michael Aizenman's class "Mathematical Physics" at Princeton in Spring 2016.

This class is of interest to both physicists and mathematicians. Several recent Fields medals are for work related to these topics.

I plan to cover the following topics. The focus is on **Topics in Mathematical Statistic Mechanics**.

1. **The statistical mechanic perspective:** systems can be described at the microscopic level with many degrees of freedom. We observe their collective behavior and find emergent behavior.
2. **Thermodynamics principles:** In departure from mechanics, which cares about equalities like $F = ma, E = mc^2$, a unique thing about thermodynamics is that its key principle is in inequality: entropy increases.

$$\Delta S \geq 0$$

3. **The emergence of thermodynamics from statistical mechanics via the equidistribution assumption and the "large deviation theory".** Mathematicians formalized the theory but the concepts were introduced earlier by physicists.
4. **Phase transitions:** A small change in parameters produces a discontinuous change in the system. For example, a small change in the temperature of water can change it from liquid to solid.
5. **Critical phenomena, critical exponents, universality classes.** Systems are macroscopically different, but the singularities you observe are given by the same power laws.
6. **Exact solution of the 2-D Ising model.** Mathematically, 3 is the hardest dimension to comprehend.
 - (a) 1-D is solvable: correlations can be described by Markov chains, and can be computed.
 - (b) In 2-D, the conformal group is very important. It gives many constraints on critical behavior, leading to a rich behavior.

- (c) In 3-D, this does not apply except for ongoing work finding consequences in 3-D from results in 2-D.
- (d) Anything with ≥ 4 dimensions gets simpler. High dimensions are characterized by the fact that loop effects do not play a large role. In 1-D a simple random walk is recurrent. In sufficient high dimensions, a simple random walk is not recurrent. The infinite-dimensional case reduces to models on trees; high dimensions exhibit the same behavior as infinite dimensions.

In coding theory and discrete math people study phase transitions of different graphs and are interested in the same topics.

I will study the 2D Ising model through “graph zeta functions.” There is a relationship between them and the solution of the graph Ising model.

7. **Stochastic geometry behind correlation functions at criticality:** in the Ising model we have a collection of spin variables $\sigma_x = \pm 1$ for $x \in \mathbb{Z}^d$. The energy favors correlations; agreement among neighbors is encouraged. These correlations spread through the system. Represent this system of correlated spins with a shadow system where we play the following game: the collection of spins is decomposed at random into connected clusters. For each such decomposition, the spins are the same in each connected cluster.

Ex. students form cliques, and each clique votes independently and unanimously. Someone who doesn’t see the cliques and just the votes sees clusters. The pattern becomes transparent once you know the cliques.

The cliques become larger and more fractal. There is an interesting fractal geometry which tells us about correlation functions.

8. **Scaling limits:** when we have a statistic mechanical system which macroscopically is described by a myriad of variables. We want to know properties of this substance. There is a technique for find these properties. This gives a link between statistical mechanics and field theory.

9. Related results for quantum spin systems.

There is a broad spectrum of references, none of which I’ll follow exclusively.

- Sacha Friedli and Yvan Velenik’s online book project (on mathematical statistical mechanics) at <http://www.unige.ch/math/folks/velenik/smbook/index.html>. I recommend this for people totally new to the subject.
- David Ruelle formulated models and basic results mathematically (late 70’s). His book helped physicists organize their thoughts. It became outdated quickly, but remains a good starting point and reference for the formalism.

Chapter 1

Introduction to statistical mechanics

2-2-16

1 The equivalence principle

1.1 Configurations and ensembles

One way to start is with the axioms of statistical mechanics. Instead I'll take a simple problem, see how it works, and present results in that context. There are simple problems that teach us a lot. The simplest is a lattice gas.

A lattice gas is a substrate where at each lattice site there may or may not be a particle.

The **configuration** is a function $n : \mathbb{Z}^d \rightarrow \{0, 1\}$:

$$n_x = \begin{cases} 1, & x \text{ is occupied,} \\ 0, & x \text{ is vacant.} \end{cases}$$

I'll use L to denote the size of the box we are considering, and $\Lambda \subset \mathbb{Z}^d$ to be a region (subset of the system). The **configuration space** is the space of possible n 's, $\{0, 1\}^\Lambda$.

We assume conservation in the number of particles, and that particles cannot overlap. We make the **equidistribution assumption**: every configuration with the same number of particles has equal probability. This gives rise to the **microcanonical ensemble**. An ensemble is a probability measure with respect to which you do averages. We have

$$\mathbb{P}(n_\Omega) = \frac{\mathbb{1}[\sum_{x \in \Omega} n_x = N]}{Z}$$

where Z is a normalization constant. Here, $\mathbb{1}[\text{cond}] := \begin{cases} 1, & \text{condition satisfied} \\ 0, & \text{else} \end{cases}$.

For every function $f : \Omega \rightarrow \mathbb{R}$ assigning a real number to each configuration, define the **microcanonical ensemble average** by

$$\langle f \rangle_{N,n}^{\text{Can}} = \frac{\sum_{n \in \Omega} \mathbb{1}[n_x = N] f(n)}{\sum_{n \in \Omega} \mathbb{1}[\sum n_x = N]}.$$

Loosely, the **equivalence principle** says that for any “local function”, the microcanonical average is approximately the grand canonical ensemble average

$$\langle f \rangle_{N,\Lambda}^{\text{Can}} \approx \langle f \rangle_{\mu,\Lambda}^{\text{Gr.C}}$$

when we take $\mu = \frac{N}{|\Lambda|}$.

The **grand canonical ensemble average** is defined as

$$\langle f \rangle_{\mu,\Lambda}^{\text{Gr.C}} = \frac{\sum_{n \in \Omega} e^{-\mu \sum_{x \in \Lambda} n_x} f(n)}{\sum_{n \in \Omega} e^{-\mu \sum_{x \in \Lambda} n_x}}$$

(Later on we will omit superscripts where it is clear.)

1.2 Equivalence principle: first proof

Consider functions which depend only on a system $\Lambda \subseteq \tilde{\Lambda}$ of much smaller volume, $|\Lambda| \ll |\tilde{\Lambda}|$. This is the sense in which the averages match up.

The micro-canonical ensemble is draconian: the number of particles is prescribed, all other configurations get weight 0. In the grand canonical ensemble, each configuration contributes. There is a value of μ where the density is the same; at that value the local average of the draconian system is asymptotically the same as that of the more relaxed system. This \approx becomes $=$ when you take the thermodynamic limit,

$$\begin{aligned} \tilde{\Lambda} &\rightarrow \mathbb{Z}^d \\ N &\rightarrow \infty \\ \frac{N}{|\tilde{\Lambda}|} &\rightarrow \rho. \end{aligned}$$

What is the induced distribution of the micro-canonical ensemble on Λ ? Under the micro-canonical ensemble what is the probability that n_Λ (the restriction to Λ) takes a particular value with $\sum_n n_x = k$? We count the number of ways to complete the configuration in $\Lambda^c = \tilde{\Lambda} \setminus \Lambda$:

$$\frac{|\{n_{\Lambda^c} : \sum_{x \in \Lambda^c} n_x = N - k\}|}{Z}$$

where Z is a normalization constant.

The number of configurations of M particles in volume V is $\binom{V}{M} = \frac{V!}{M!(V-M)!}$. Using Stirling's approximation

$$\ln(M!) = M(\ln M - 1)(1 + o(1)),$$

letting the **entropy** S be the logarithm of the number of configurations and

$$s(\rho) = -[\rho \ln \rho + (1 - \rho) \ln(1 - \rho)],$$

we have

$$\binom{V}{M} = \frac{V!}{M!(V-M)!} =: e^{S(M,V)} \approx e^{Vs(\rho)}$$

(do this calculation as an exercise).

Shannon also found such a formula for entropy.

This attains maximum of $\ln 2$ at $\frac{1}{2}$ where it has quadratic behavior.

The implication is that if you slightly change the density, the number of configurations changes drastically. In physical substances V may be 10^{23} . The change would then be $e^{10^{23}\Delta s}$. In any average over configurations, only those at the peak contribute: “winner takes all.”

What is the probability of observing k particles in the small box given n in the big box?

$$\begin{aligned} \mathbb{P}(n_\Lambda) &\approx \frac{e^{|\Lambda^c|s\left(\frac{N-k}{|\widetilde{\Lambda}|-|\Lambda|}\right)}}{C} \\ \frac{N-k}{|\widetilde{\Lambda}|-|\Lambda|} &= \rho - \underbrace{\left(\rho - \frac{N}{|\widetilde{\Lambda}|-|\Lambda|}\right)}_{\rightarrow 0 \text{ independent of } k} - \frac{k}{|\widetilde{\Lambda}|-|\Lambda|} \\ s\left(\frac{N-k}{|\widetilde{\Lambda}|-|\Lambda|}\right) &\approx s(\rho) - s'(\rho)\frac{k}{|\Lambda^c|} \end{aligned}$$

Changing k by a little bit affects how many particles are outside but not so much the density outside: the correction term is small. Hence for $\sum_{x \in \Lambda} n_x$,

$$\mathbb{P}(n_\Lambda) \approx \frac{e^{|\Lambda^c|s(\rho)} e^{-s'(\rho)k}}{Z}.$$

$e^{|\Lambda^c|s(\rho)}$ is a huge factor but it does not vary with k so we can omit it. Writing out Z as the sum of the numerators over all n_Λ , and letting $\mu = s'(\rho)$, we get that this equals

$$= \frac{e^{-\mu k}}{\sum_{n' \in \Omega_\Lambda} e^{-\mu \sum_{x \in \Lambda} n'_x}},$$

which is the grand canonical ensemble average. The rest of the system acts on the small system as a “particle (heat) bath.”

Here we used very explicit machinery, namely, the Stirling formula. We want a general expression that doesn't rely on the Stirling formula because in more complicated models, we will not have the luxury of using the Stirling approximation. We'll do this in the next section.

Then we will be able to consider models where there are more energy constraints. The general procedure is first make a list of energy constraint. There is a generalization of the equivalence principle where

- in the micro-canonical ensemble we average over configurations where the constraints have prescribed values. (We considered the special case where just the number of particles was prescribed.)

- in the grand canonical ensemble, we add an energy term to the exponential: $e^{-\mu N(n) - \beta \mathcal{E}(n)}$ where $N(n)$ is the number of particles, $-\mu N(n)$ is the Gibbs factor, and $\mathcal{E}(n)$ the energy.

How can we construct an alternative method without the Stirling formula?

Define

$$Z_{\tilde{\Lambda}} = \sum_{n \in \Omega} e^{-\mu N(n)}.$$

This can be easily computed without Stirling. From the value of this, you can learn the value of the entropy function:

$$\begin{aligned} &= \sum_{n \in \Omega} \prod_{x \in \tilde{\Lambda}} e^{-\mu \mathbb{1}[n_x=1]} \\ &= \prod_{x \in \tilde{\Lambda}} (1 + e^{-\mu}) = (1 + e^{-\mu})^{|\tilde{\Lambda}|}. \end{aligned}$$

However,

$$\begin{aligned} Z_{\tilde{\Lambda}} &= \sum_{n \in \Omega} e^{-\mu N(n)} \\ &= \sum_{K \in \mathbb{N}} e^{-\mu K} e^{Vs(\frac{K}{V})} \\ &= \sum_{K \in \mathbb{N}} e^{V(s(\frac{K}{V}) - \mu \frac{K}{V})}. \end{aligned}$$

For each k count how many configurations have that value of k . Here $S(V, K) = Vs(\frac{K}{V})$. Suppose we find it acceptable to say the system *has* an entropy.

The maximal value of $\frac{K}{V}$ takes it all: this is

$$= e^{V \max_{\rho \in [0,1]} [s(\rho) - \mu \rho]}.$$

This is the **Legendre transform** of the entropy. Let

$$\text{eq:s-leg} s^*(\mu) := \max_{\rho} [s(\rho) - \mu \rho] = \ln(1 + e^{-\mu}) \quad (1.1)$$

Next time we'll discuss how to derive from this expression the formula for $s(\rho)$, using the inverse Legendre transform.

“The elementary problems are the most precious, once you absorb them they are part of your makeup.”

2-4-16

1.3 Second proof

The **partition function** for the grand canonical ensemble is

$$\begin{aligned} Z_{\Lambda} &= \sum_{n \in \Omega} e^{-\mu N(n)} = (1 + e^{-\mu})^{|\Lambda|} \\ \frac{1}{|\Lambda|} \ln Z_{\Lambda} &= 1 + e^{-\mu} \end{aligned}$$

At each site there are 2 possibilities, and the different sites are independent. We find¹

$$\lim_{|\Lambda| \rightarrow \infty} \frac{1}{|\Lambda|} \ln Z_\Lambda = \sup_{0 \leq \rho \leq 1} [s(\rho) - \rho\mu]. \quad (1.2)$$

In slow motion, note that the only thing that matters in the summand is the number of particles in n , so let's group the summands by this. Letting $\rho = \frac{N}{|\Lambda|}$,

$$\begin{aligned} Z_N &= \sum_{n \in \Omega} e^{-\mu N(n)} \\ &= \sum_{n \in \Omega} e^{-\mu N(n)} e^{|\Lambda| s(\rho)} \\ &= \sum_{\rho} e^{|\Lambda| [s(\rho) - \mu\rho]}. \end{aligned}$$

The multiplicity is the exponential of the volume times the entropy. The error we are making by focusing on the maximal point rather than counting with multiplicity the near-maximal value. We lose at most a factor equal to the volume. But $\frac{1}{|\Lambda|} \ln |\Lambda| \rightarrow 0$.

We find

$$\max_{\rho} [s(\rho) - \mu\rho] \leq \frac{1}{|\Lambda|} \ln Z_\Lambda \leq \max_{\rho} [s(\rho) - \mu\rho] + \underbrace{\frac{1}{|\Lambda|} \ln |\Lambda|}_{\rightarrow 0}$$

Hence

$$\sum_{0 \leq \rho \leq 1} [s(\rho) - \mu\rho] = 1 + e^{-\mu}$$

For $s(\rho) = -[\rho \ln \rho + (1 - \rho) \ln(1 - \rho)]$, we find the critical point by setting the derivative to 0. Trick: $[x(\ln x - 1)]' = \ln x$. We can subtract 1 from each of the logs changing the expression by a constant. Thus

$$s'(\rho) = -\ln \rho + \ln(1 - \rho) - \mu.$$

Solving gives $\frac{\rho}{1-\rho} = e^{-\mu}$.

1.4 Legendre transform

In general, a micro-canonical ensemble specifies all the conserved quantities: Particles, energies, and whatever else is there. The grand canonical ensemble also generalizes by changing the factor to $e^{-\mu N} e^{\beta H}$. Such factors are referred to as **Gibbs factors** or Gibbs measures. Next topics.

- Legendre transform
- Convexity/concavity

¹Mathematicians are paranoid, so we use sup instead of max. For many practical purposes they are the same.

- A variational characterization of Gibbs states
- First order phase transitions (thermodynamics and statistical mechanics perspectives)

Definition 1.1.1: A function on \mathbb{R}^k is **concave** (**convex**) if for any $x_0, x_1 \in \mathbb{R}^k$, $0 \leq \lambda \leq 1$,

$$F(\lambda x_1 + (1 - \lambda)x_0) \begin{matrix} \leq \\ \geq \end{matrix} \lambda F(x_1) + (1 - \lambda)F(x_0).$$

Concavity (convexity) means if you draw a chord between two points, it will lie below (above) the curve.

For a strictly concave function, a maximum, whenever it exists, is unique.

Theorem 1.1.2: For any concave function on \mathbb{R} ,

1. The directional derivatives $F'(x \pm 0)$ exist at all $x \in \mathbb{R}$. The **directional derivative** is defined as


$$F'(x + 0) = \lim_{\varepsilon \rightarrow 0^+} \frac{F(x + \varepsilon) - F(x)}{\varepsilon}$$

$$F'(x - 0) = \lim_{\varepsilon \rightarrow 0^-} \frac{F(x + \varepsilon) - F(x)}{\varepsilon}.$$

2. $F'(x - 0) \geq F'(x + 0)$ and $F'(x \pm 0)$ are decreasing.
3. For all but countably many values $x \in \mathbb{R}$, $F'(x - 0) = F'(x + 0)$, i.e., F' 's different at x .
4. Let F_n be a sequence of concave functions which converge pointwise: for all x , $\lim_{N \rightarrow \infty} F_N(x) =: \tilde{F}(x)$ exists. Then
 - (a) \tilde{F} is concave.
 - (b) At points of differentiability of \tilde{F} , the derivatives also converge,


$$F'(x \pm 0) \rightarrow \tilde{F}'(x).$$

²Finite energy functions are always smooth, but their limit can have discontinuous derivative.



Much of this generalizes to directional derivatives in n dimensions.


Proof. Concavity implies that the slope of the line between $x, x + \varepsilon$ is increasing as $\varepsilon \rightarrow 0^+$.



To prove 3, think of the graph of the derivative at the right. The sum of any uncountable number of steps is infinite.^{3 4}

³More rigorously, for every nonzero interval $[F'(x + 0), F'(x - 0)]$, we can associate with it a rational number. \mathbb{Q} is countable

⁴The set of discontinuities can be dense, ex. at all the rational numbers. In physics it was thought that this can't occur, but there are materials whose free energy discontinuities is dense in certain areas.



□


Definition 1.1.3: The **Legendre transform** of a function is defined as

$$(TG)(y) = \inf_x [y \cdot x - G(x)] = -\sup_x [G(x) - y \cdot x]$$

Theorem 1.1.4: For any function G , TG is concave.

Proof. An efficient way to think about this: for each value of the parameter x , as a function of y this is a linear function. For each x we get a linear function. Define the transform by taking the infimum over that.

Take 2 points and draw the chord between them. For each linear function the chord lies below it.




□

I.e., we used

Proposition 1.1.5: Let \mathcal{F} be a collection of concave (e.g. linear) functions. Then $\inf_{f \in \mathcal{F}} f(x)$ is concave.

Definition 1.1.6: The **concave hull** \tilde{G} of G is defined as the smallest concave function that is at least the function value at every point:

$$\tilde{G}(x) = \inf \{F(x) : F \text{ concave}, \forall u, F(u) \geq G(u)\}.$$



Theorem 1.1.7: For *concave* G ,

$$T(TG) = G.$$

In general, $T(TG)$ is the concave hull of G .

Proof for G differentiable. Use the fact that if G is differentiable, then

$$\inf_x [y \cdot x - G(x)]$$

occurs at $y = G'(x)$. □

Note that we plotted the function and the dual function on the same graph. However, they have inverse units, for example, energy and inverse temperature. You get a lot of insights into physics if you keep track of the units.

Recall $\rho = \frac{N}{|\Lambda|}$. If all $2^{|\Lambda|}$ configurations are given equal weight, the typical value of ρ , the particles per unit volume is $\frac{1}{2}$. The Law of Large Numbers says that with probability 1 the ratio tends to $\frac{1}{2}$. Is it possible that the density is $\frac{1}{3}$? Yes, but the probability of such a density is given by the entropy: it's exponentially small, $e^{|\Lambda|[s(\frac{1}{3}) - \ln 2]}$. Anything other than $\frac{1}{2}$ is a large deviation; they occur with exponentially small probability. We want to quantify the probability of large deviation events. Here is a language that people found useful.

Definition 1.1.8: A sequence of probability measures on \mathbb{R} is said to satisfy a **large deviation principle** with speed $\{a_N\}$, and rate function I if for each $x \in \mathbb{R}$, $\varepsilon > 0$,

$$-\inf_{|u-x|<\varepsilon} I(u) \leq \liminf_{N \rightarrow \infty} \frac{1}{a_N} \ln \mathbb{P}_N((x-\varepsilon, x+\varepsilon]) \leq \limsup_{N \rightarrow \infty} \frac{1}{a_N} \ln \mathbb{P}_N((x-\varepsilon, x+\varepsilon]) \leq -\inf_{|x-u| \leq \varepsilon} I(u)$$

For us, $I(x) = \ln 2 - s(x)$.

Index

concave, 10
concave hull, 13
configuration, 5
convex, 10

directional derivative, 10

entropy, 6
equidistribution assumption, 5
equivalence principle, 6

Gibbs factors, 9
grand canonical ensemble average, 6

large deviation principle, 13
Legendre transform, 8, 12

microcanonical ensemble, 5
microcanonical ensemble average, 5

partition function, 8